Computer Control Networks

Q4

Internetwork Quality of Service

timing

less than 1ms

Transport Protocols

- Internet traffic continues to grow and change
  - demand for real-time responses
  - increasing use of audio, image, and video
  - heavy use of the World Wide Web
- packet-switching technology with routers functioning as switches was not designed to handle voice and video
- strong need to support a variety of traffic with QoS requirements within TCP/IP

Integrated Service Architecture (ISA)

classify the data

- intended to provide QoS transport over IP-based Internets
- defined in RFC 1633
- portions already being implemented in routers and end-system software

Internet Traffic - Elastic

- traffic that can adjust, over wide ranges, to changes in delay and throughput and still meet the needs of its applications
- traditional type of traffic supported on TCP/IP-based Internets
- applications classified as elastic include:

timing

elastic

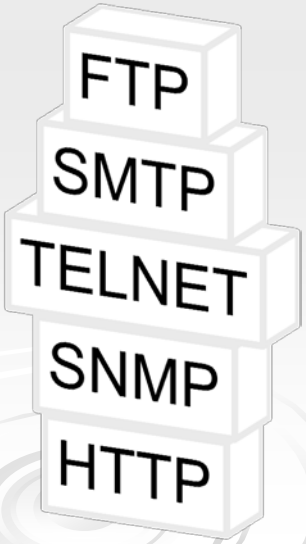

Internet Traffic - Inelastic

- does not easily adapt, if at all, to changes in delay and throughput across an internet
  - real-time traffic

requirements for inelastic traffic include:

- Throughput
- Delay
- Jitter
- Packet loss

- new internet architecture requirements:
  - resource reservation protocol
  - elastic traffic still needs to be supported

only for small network.
need to know the whole network

ISA

- purpose is to enable QoS support over IP-based internets
- sharing capacity during congestion is the central design issue

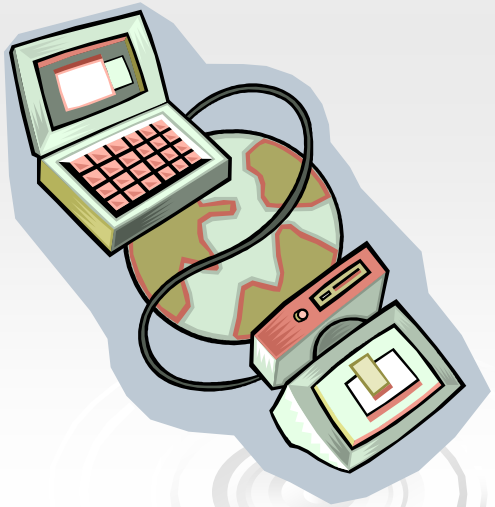

- to manage congestion and provide QoS transport ISA makes use of:
  - Admission control
  - Routing algorithm
  - Queuing discipline
  - Discard policy

Integrated Service Architecture.

ISA Components

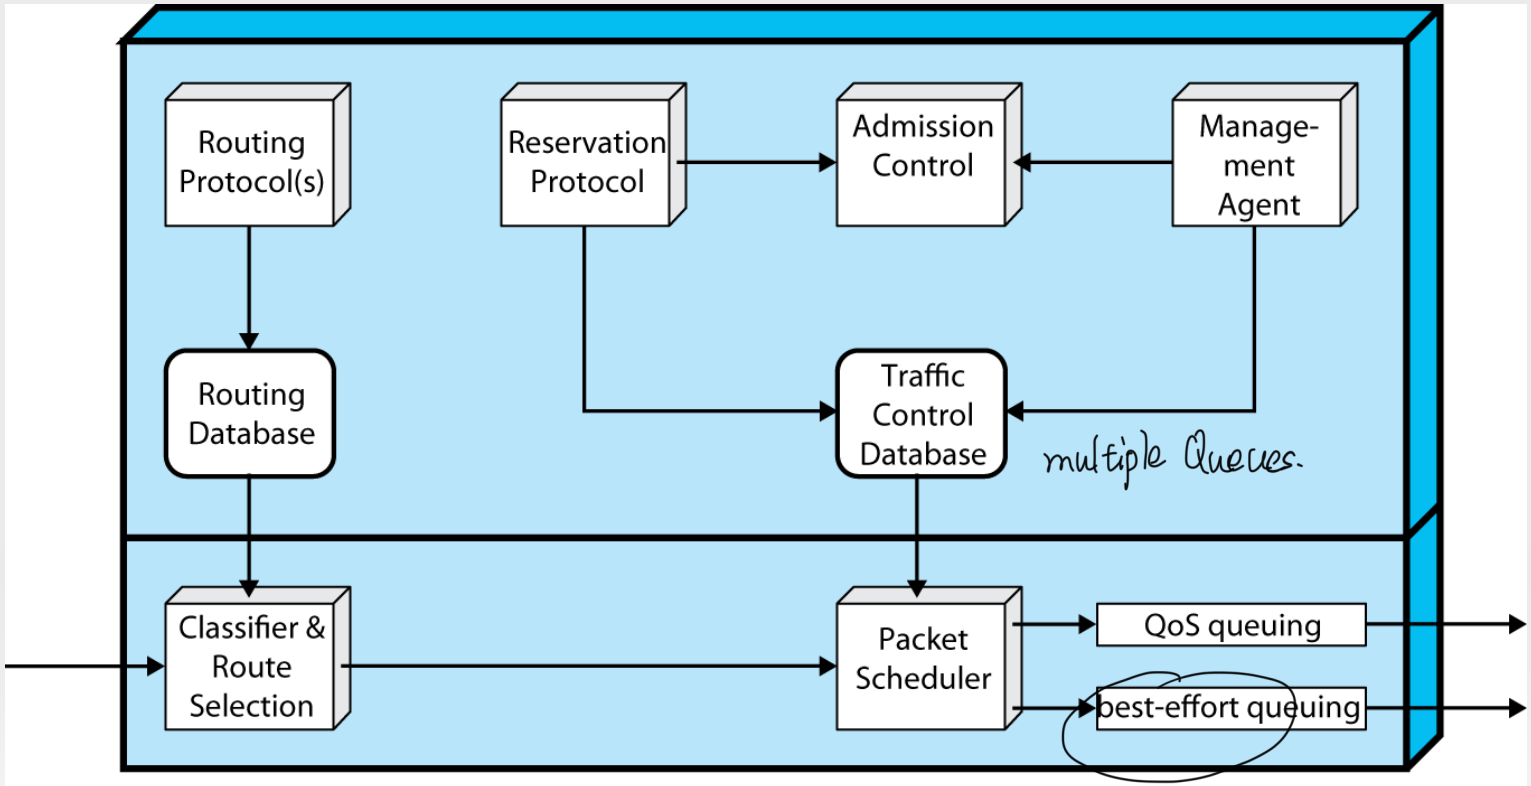

Figure 20.1 Integrated Services Architecture Implemented in Router

ISA Services

- traffic specification (TSpec)

three categories of service:

- Guaranteed
- Controlled load
- Best effort

token bucket traffic scheme

- traffic sources can be defined easily and accurately
- provides a concise description of the load
- provides the input parameters to a policing function

Token Bucket Scheme

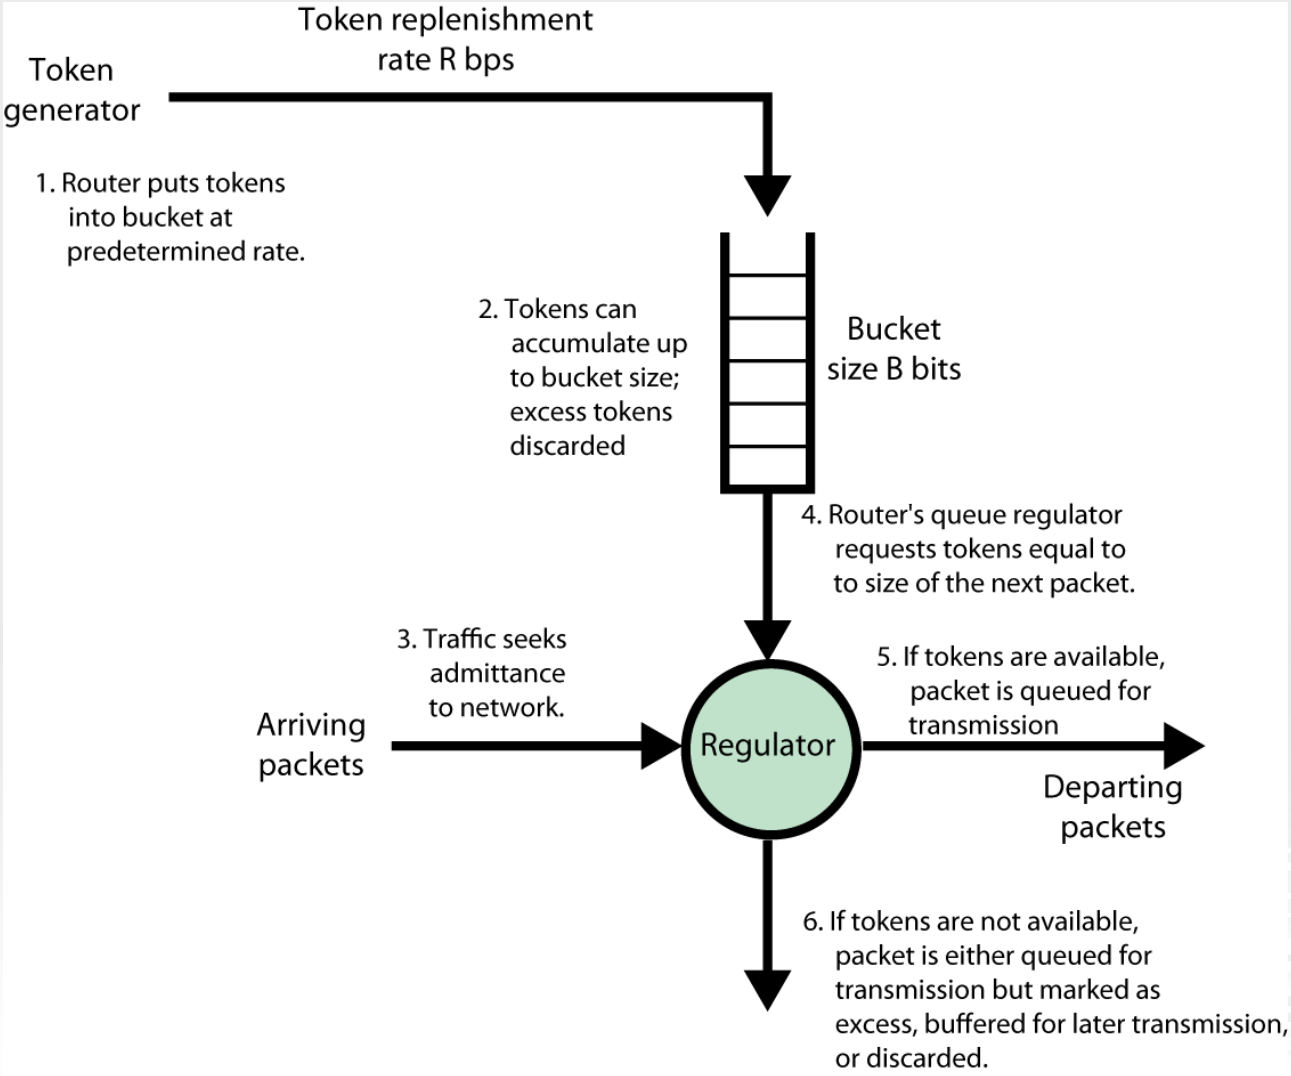

Figure 20.2 Token Bucket Scheme

Guaranteed Service

- key elements are:
  - service provides assured capacity
  - specified upper bound on the queuing delay through the network
  - there are no queuing losses
- application provides a characterization of expected traffic profile and the service determines the end-to-end delay that it can guarantee
- most demanding service provided by ISA

Controlled Load

- key elements are:
  - tightly approximates the behavior visible to applications receiving best-effort service under unloaded conditions
  - no specified upper bound on the queuing delay through the network
  - high percentage of transmitted packets will be successfully delivered
- useful for adaptive real-time applications

Queuing Discipline

- routers use first-in-first-out (FIFO) queuing discipline

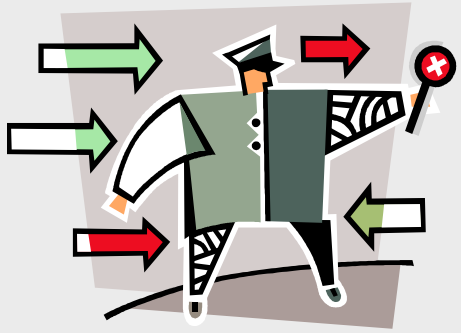

- drawbacks of FIFO
  - no special treatment given to higher priority packets
  - smaller packets get delayed behind larger packets
  - a greedy TCP connection can crowd out more altruistic connections

FIFO and Fair Queuing

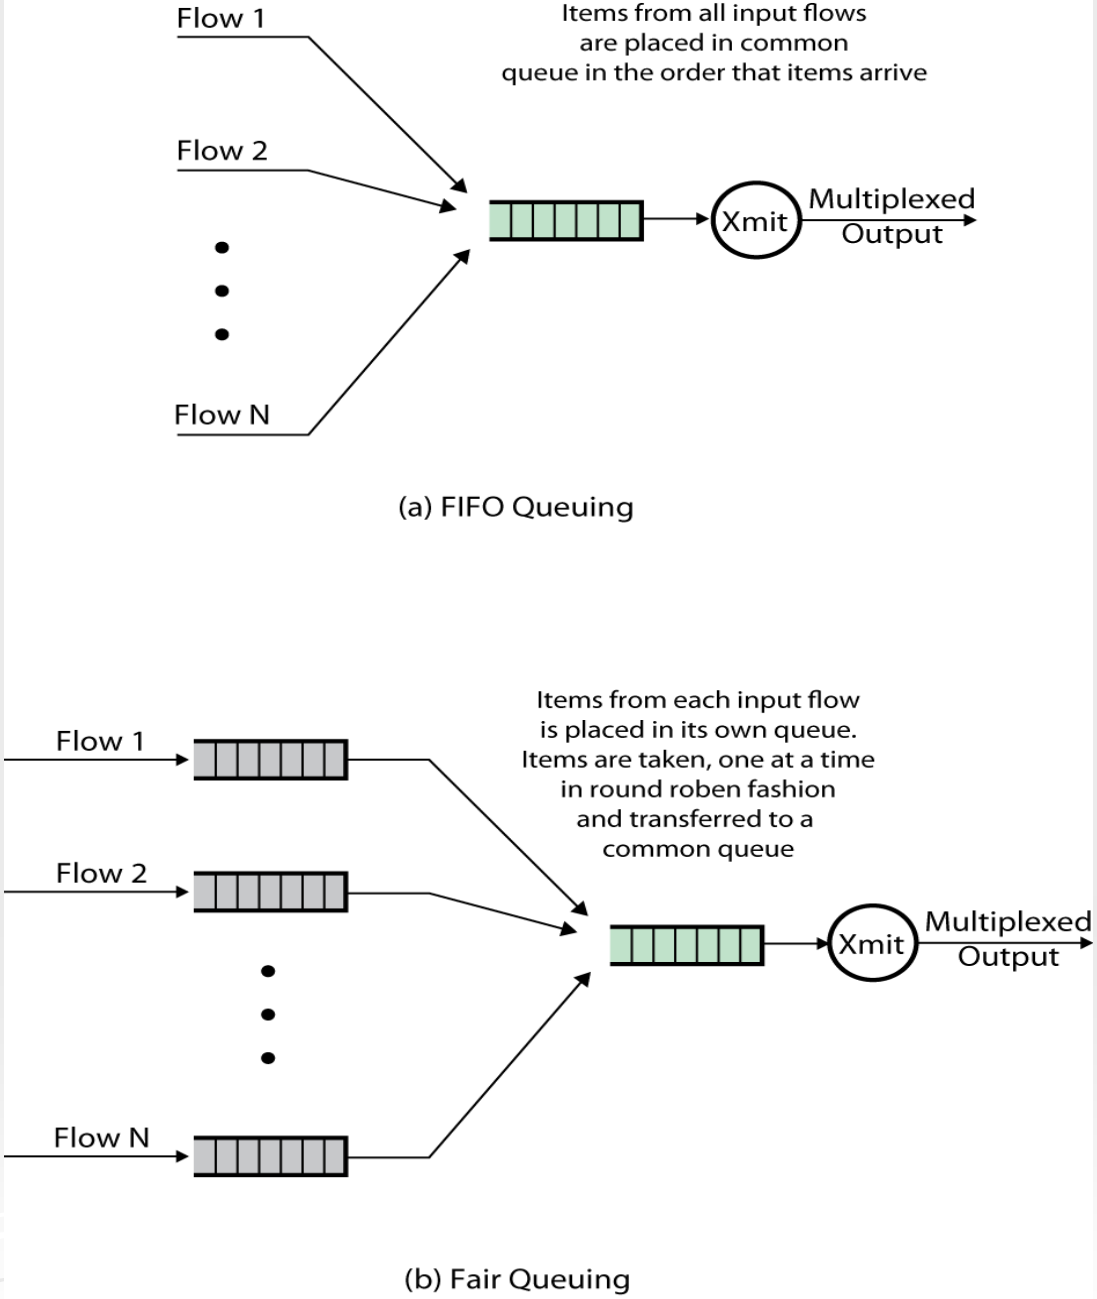

Figure 20.3 FIFO and Fair Queuing

Resource ReSerVation Protocol (RSVP)

- RFC 2205
- provides supporting functionality for ISA
- prevention strategy
  - have unicast applications reserve resources in order to meet a given QoS
  - enables routers to decide ahead of time if they can meet the delivery requirement for a multicast transmission
- must interact with a dynamic routing strategy
  - soft state

RSVP Characteristics

- *Unicast and multicast*
- Simplex
- *Receiver-initiated reservation*
- Maintaining soft state in the internet
- *Providing different reservation styles*
- Transparent operation through non-RSVP routers

Receiver-Initiated Reservation

- since receivers specify the desired QoS it makes sense for them to make resource reservations
  - different members of the same multicast group may have different resource requirements
  - QoS requirements may differ depending on the output equipment, processing power, and link speed of the receiver
  - routers can aggregate multicast resource reservations to take advantage of shared path segments

Soft State

- connectionless
- reservation state is cached information in the routers that is installed and periodically refreshed
- if a new route becomes preferred the end systems provide the reservation to the new routers on the route

Data Flows

- **Basis of RSVP operation:**

Session

- Destination IP address
- IP protocol identifier
- Destination port

Flow specification

- Service class
- Rspec
- Tspec

Filter specification

- Source address
- UDP/TCP source port

Data Flow Relationship

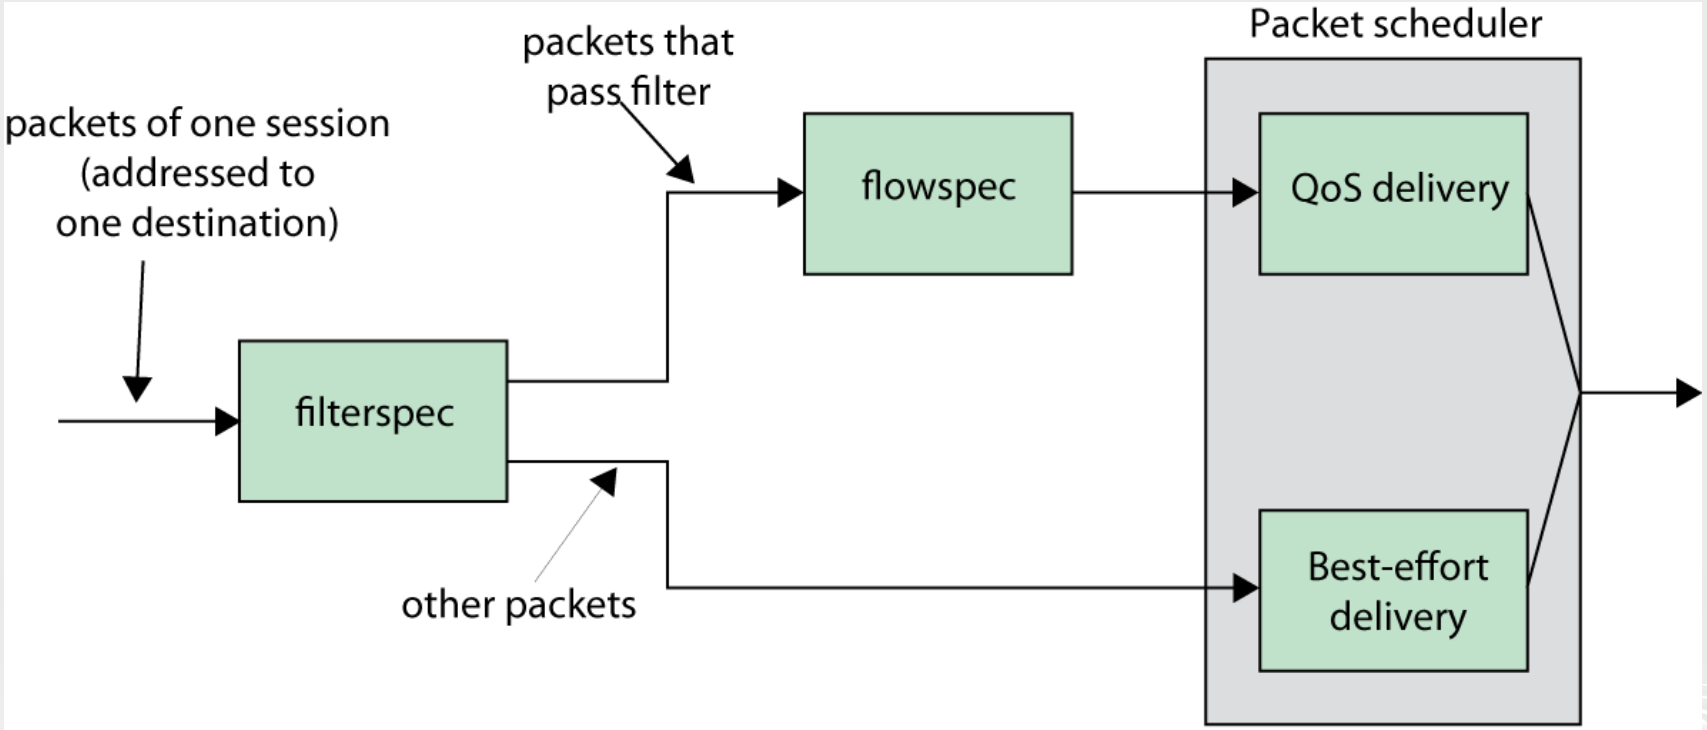

Figure 20.4 Treatment of Packets of One Session at One Router

Differentiated Services (DS)

- RFC 2475
- designed to provide a tool to support a range of network services
- key characteristics:
  - no change to IP is required
  - SLA is established prior to use of DS
    - applications do not need to be modified
  - provides a built-in aggregation mechanism
    - good scaling to larger networks and traffic loads
  - DS is implemented in individual routers
  - most widely accepted QoS in enterprise networks

Terminology for Differentiated Services

Behavior Aggregate	A set of packets with the same DS codepoint crossing a link in a particular direction.
Classifier	Selects packets based on the DS field (BA classifier) or on multiple fields within the packet header (MF classifier).
DS Boundary Node	A DS node that connects one DS domain to a node in another domain
DS Codepoint	A specified value of the 6-bit DSCP portion of the 8-bit DS field in the IP header.
DS Domain	A contiguous (connected) set of nodes, capable of implementing differentiated services, that operate with a common set of service provisioning policies and per-hop behavior definitions.
DS Interior Node	A DS node that is not a DS boundary node.
DS Node	A node that supports differentiated services. Typically, a DS node is a router. A host system that provides differentiated services for applications in the host is also a DS node.
Dropping	The process of discarding packets based on specified rules; also called **policing**.
Marking	The process of setting the DS codepoint in a packet. Packets may be marked on initiation and may be re-marked by an en route DS node.
Metering	The process of measuring the temporal properties (e.g., rate) of a packet stream selected by a classifier. The instantaneous state of that process may affect marking, shaping, and dropping functions.
Per-Hop Behavior (PHB)	The externally observable forwarding behavior applied at a node to a behavior aggregate.
Service Level Agreement (SLA)	A service contract between a customer and a service provider that specifies the forwarding service a customer should receive.
Shaping	The process of delaying packets within a packet stream to cause it to conform to some defined traffic profile.
Traffic Conditioning	Control functions performed to enforce rules specified in a TCA, including metering, marking, shaping, and dropping.
Traffic Conditioning Agreement (TCA)	An agreement specifying classifying rules and traffic conditioning rules that are to apply to packets selected by the classifier.

DS Services

- typically DS domain is under the control of one administrative entity
- services provided across a DS domain are defined in an SLA

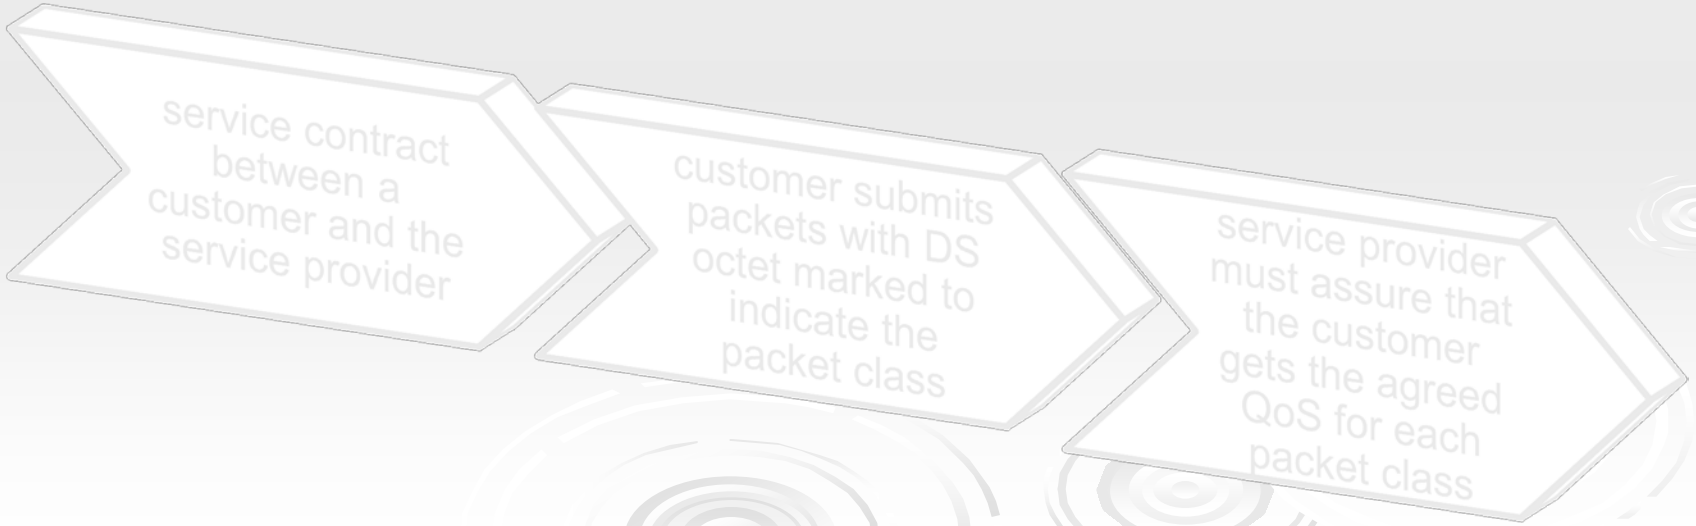

Performance Parameters in an SLA

- detailed service performance parameters such as expected throughput, drop probability, latency
- constraints on the ingress and egress points at which the service is provided, indicating the scope of the service
- traffic profiles that must be adhered to for the requested service to be provided, such as token bucket parameters
- disposition of traffic submitted in excess of the specified profile

Services Provided in a SLA

- traffic offered at service level A will be delivered with low latency
- traffic offered at service level B will be delivered with low loss
- ninety percent of in-profile traffic delivered at service level D will be delivered
- traffic offered at service level E will be allotted twice the bandwidth of traffic delivered at service level F
- traffic with drop precedence X has a higher probability of delivery than traffic with drop precedence Y

DS Field

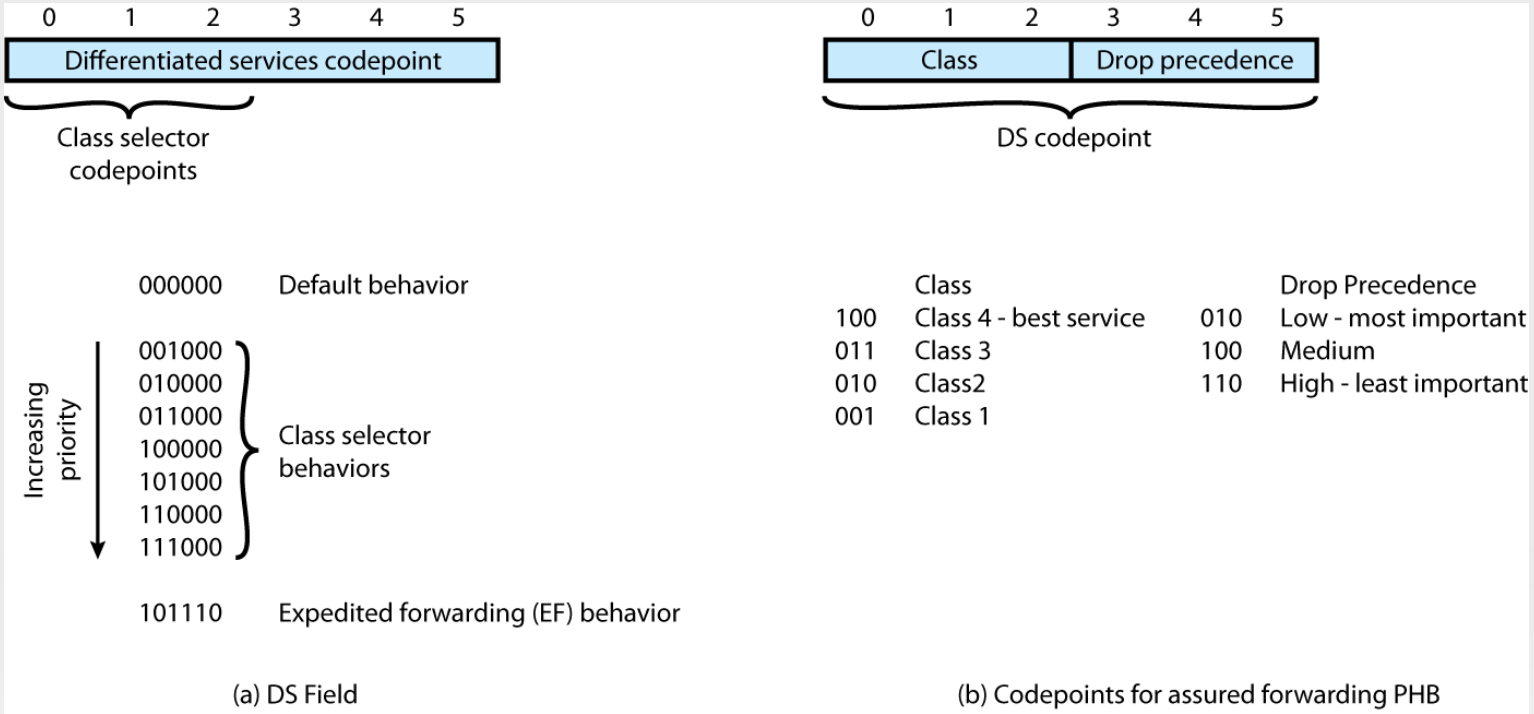

Figure 20.8 DS Field

DS Domains

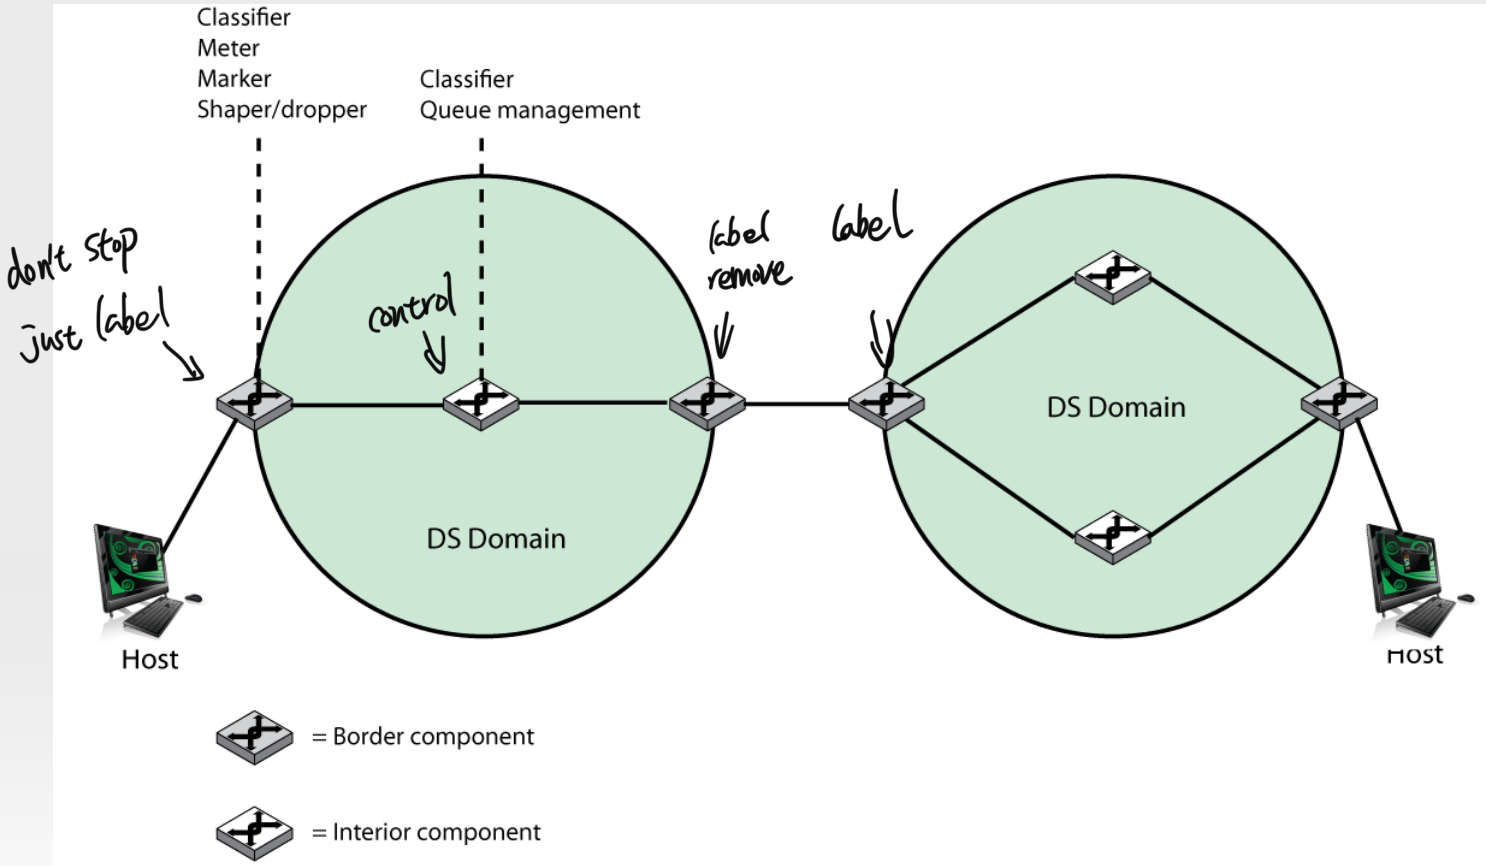

Figure 20.9 DS Domains

DS Traffic Conditioner

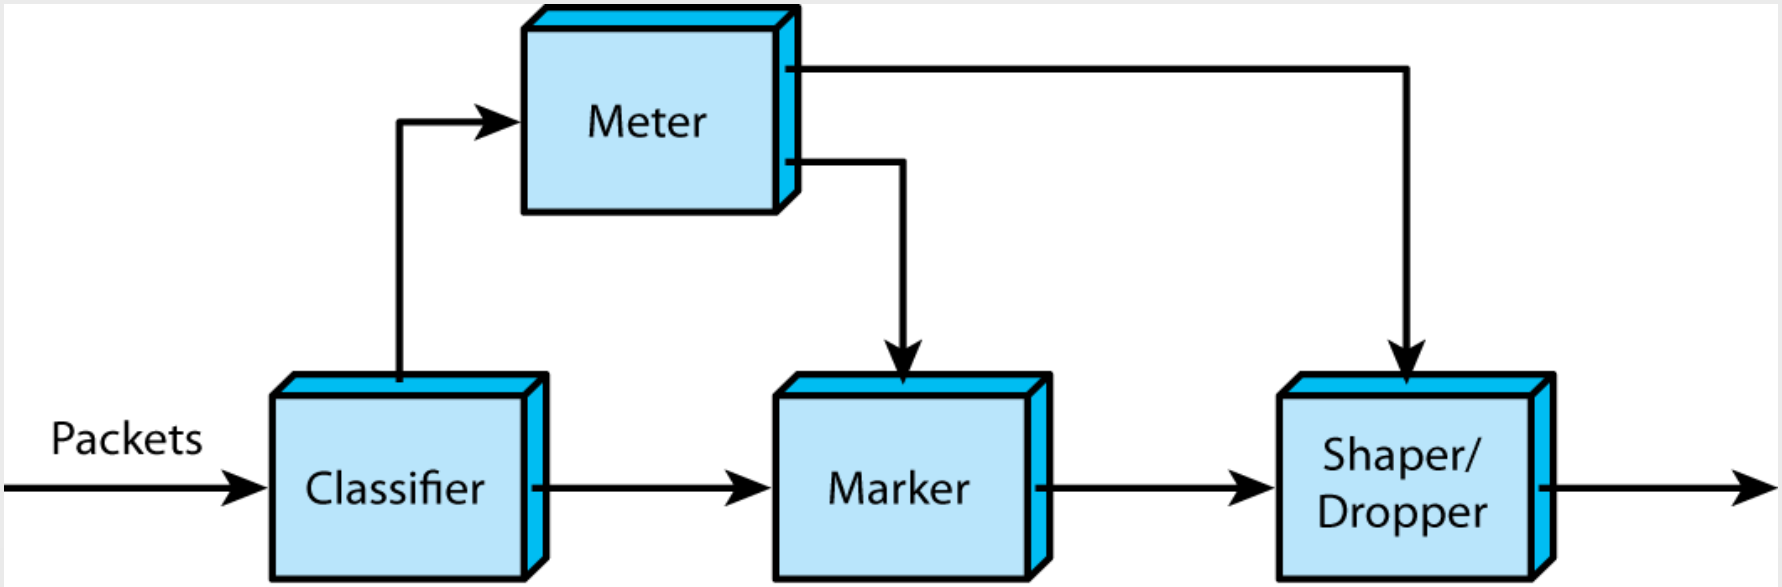

Figure 20.10 DS Traffic Conditioner

Expedited Forwarding PHB (EF PHB)

- RFC 3246
- building block for low-loss, low-delay, and low-jitter end-to-end services through DS domains
  - difficult to achieve
  - cause is queuing behavior at each node
  - intent is to provide a PHB in which packets encounter short or empty queues
  - configures nodes so traffic has a well-defined minimum departure rate

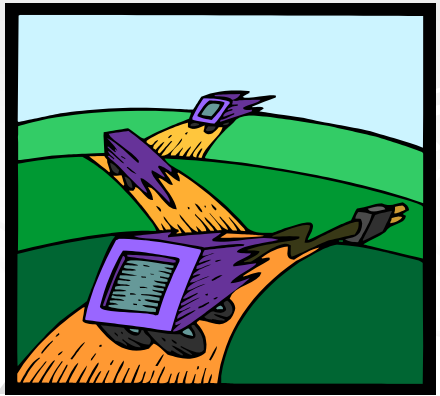

Assured Forwarding PHB (AF) PHB

- RFC 2597
- designed to provide a service superior to best-effort but one that does not require the reservation of resources within an internet
- referred as to as explicit allocation
  - expands by defining four AF classes and marking packets with one of three drop precedence values

Service Level Agreements (SLA)

- contract between a network provider and a customer that defines specific aspects of the service that is to be provided

SLA includes:

- a description of the nature of service to be provided
- the expected performance level of the service
- the process for monitoring and reporting the service level

Framework for SLA

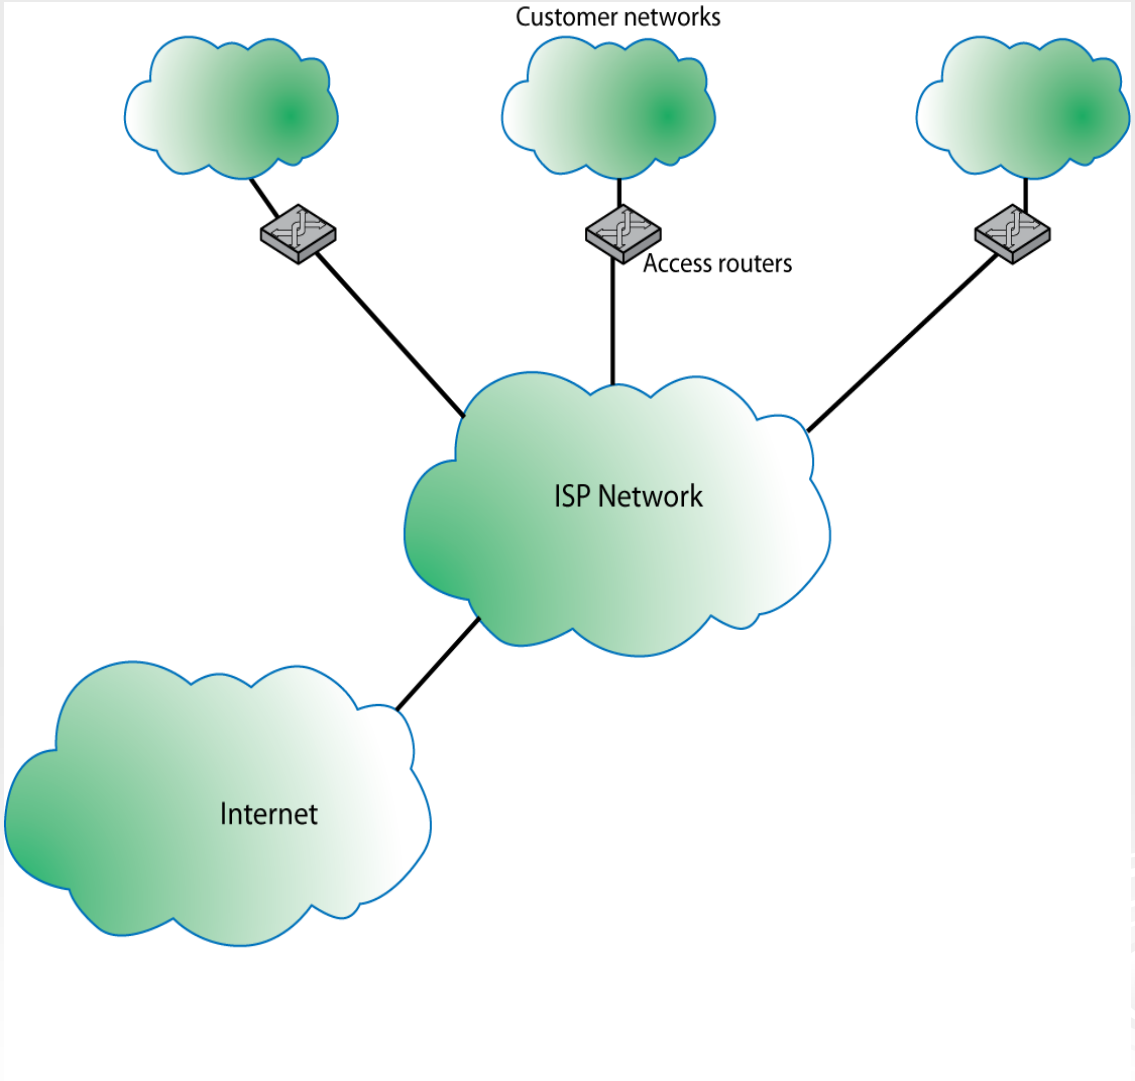

Figure 20.11 Typical Framework for Service Level Agreement

loss & delay

IP Performance Metrics

- IPPM chartered by IETF to develop standard metrics that relate to the quality, performance, and reliability of Internet data delivery
- need for standardization:
  - Internet has grown and continues to grow at a dramatic rate
  - Internet serves a large and growing number of commercial and personal users across an expanding spectrum of applications

loss & delay

normal voice 2% loss.

IP Performance Metrics

no loss

no delay

client -> server 4ms

failure

Minimum packet

(a) Sampled metrics

Metric Name	Singleton Definition	Statistical Definitions
One-Way Delay	Delay = dT, where Src transmits first bit of packet at T and Dst received last bit of packet at T + dT	Percentile, median, minimum, inverse percentile
Round-Trip Delay	Delay = dT, where Src transmits first bit of packet at T and Src received last bit of packet immediately returned by Dst at T + dT	Percentile, median, minimum, inverse percentile
One-Way Loss	Packet loss = 0 (signifying successful transmission and reception of packet); = 1 (signifying packet loss)	Average
One-Way Loss Pattern	Loss distance: Pattern showing the distance between successive packet losses in terms of the sequence of packets Loss period: Pattern showing the number of bursty losses (losses involving consecutive packets)	Number or rate of loss distances below a defined threshold, number of loss periods, pattern of period lengths, pattern of inter-loss period lengths.
Packet Delay Variation	Packet delay variation (pdv) for a pair of packets with a stream of packets = difference between the one-way-delay of the selected packets	Percentile, inverse percentile, jitter, peak-to-peak pdv

Src = IP address of a host
Dst = IP address of a host

(b) Other metrics

Metric Name	General Definition	Metrics
Connectivity	Ability to deliver a packet over a transport connection.	One-way instantaneous connectivity, Two-way instantaneous connectivity, one-way interval connectivity, two-way interval connectivity, two-way temporal connectivity
Bulk Transfer Capacity	Long-term average data rate (bps) over a single congestion-aware transport connection.	BTC = (data sent)/(elapsed time)

Packet Delay Variation

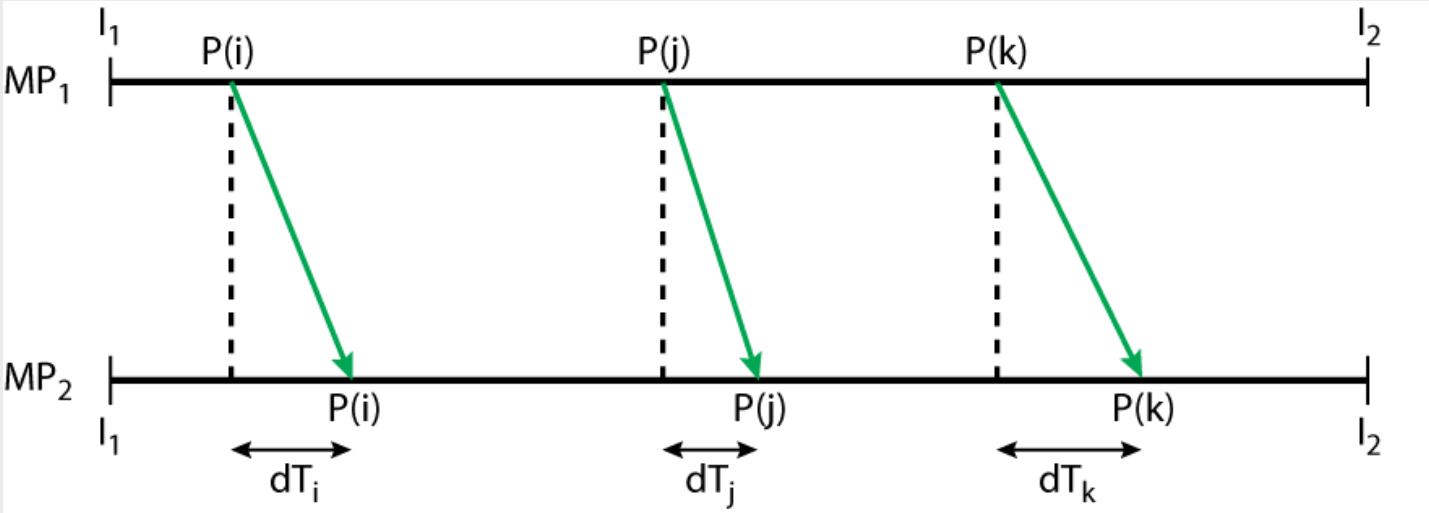

I_1, I_2 = times that mark that beginning and ending of the interval in which the packet stream from which the singleton measurement is taken occurs.

MP_1, MP_2 = source and destination measurement points

P(i) = ith measured packet in a stream of packets

dT_i = one-way delay for P(i)

Figure 20.12 Model for Defining Packet Delay Variation

Summary

- Integrated services architecture
  - Elastic and inelastic traffic
  - ISA approach, components, services
- Resource reservation protocol
  - RSVP goals and characteristics
  - RSVP operation, filtering, reservation styles
- Differentiated services
  - DS fields, configuration, operation
- Service level agreements
- IP performance metrics